

Clinical Scope – User Guide

Clinical Scope Team

September 7, 2026

Contents

1	Introduction	3
1.1	Key Features	3
2	Launching the Application	4
2.1	First launch	4
2.2	Starting the App	4
2.3	Application Overview	5
3	Patient Data & Supported Data Sources	7
3.1	Example Folder Layouts	7
3.2	Folder Naming Rules	7
3.3	Canonical Data Source Table	7
3.4	Generic “Other” Data Source	9
4	Loading Database Options	11
4.1	Option 1: Upload a Custom Configuration File	11
4.2	Option 2: Reload Last Config (Daily Workflow)	11
4.3	Option 3: Default Visualization (Quick Start)	11
5	Configuring Patient Options	12
5.1	Global Options	12
5.2	Per-Source Options	12
6	Settings	14
6.1	App behavior	14
6.2	Plot defaults	14
7	Processing Data	16
7.1	Process Visualization (orange button)	16
7.2	Inspect Data (teal button)	16
7.3	Inspecting only the signals you configured	17
7.4	Inspecting from the Command Line	18
8	Interacting with Plots	19
8.1	Navigation Controls	19
8.2	Dynamic Resampling (FigureResampler)	19
9	Annotations	20
9.1	Annotation Toolbar	20
9.2	Annotation Types	20
9.3	Moving an Annotation	20

9.4	Dragging an Annotation with the Mouse	20
9.5	Hiding an Annotation	21
9.6	Groups	21
9.7	Persistence	21
9.8	Python API	21
10	Configuration File Reference	23
10.1	patient_options.json	23
10.2	database_options.json	24
10.3	database_options.xlsx	34
11	Troubleshooting	40
11.1	Browser Does Not Open Automatically	40
11.2	No Data Found	40
11.3	Slow Loading	40
11.4	Time Alignment Issues	40
11.5	Application Crashes or Errors	40
11.6	Known limitations	40

1 Introduction

Clinical Scope is an interactive dashboard for exploring, comparing, and annotating clinical physiological signals from multiple medical devices in a single interface.

1.1 Key Features

- **Multi-source visualization:** Display signals from FluxMed, Mindray, EIT, Servo-U and EDF recorders simultaneously.
- **Generic “Other” data source:** Drop any CSV or Parquet file with a datetime column into an **other/** folder — monitor exports, syringe pumps, anything tabular. Signals are auto-discovered and configured per file.
- **Interactive plots:** Zoom, pan, and explore data at any time scale with automatic resampling.
- **Data inspection:** Preview available columns, point counts, and time ranges for every source before running the full visualization — exportable to CSV.
- **Annotations:** Draw lines and rectangles directly on plots to mark events or regions of interest. Annotations are saved and persist across sessions.
- **Flexible configuration:** Choose which signals to display, customize labels, units, colors, and group related signals together — in JSON or Excel format.
- **Cross-datasource phase loops:** Define loop entries to build phase plots that combine signals from different devices.
- **Spectrograms:** Render a time-vs-frequency-vs-power view of any signal (e.g. EEG) over a configurable frequency band, with a fixed colour range for consistent reading across patients.
- **Power spectral densities (PSD):** Render power against frequency, averaged over the loaded time range, with several signals overlaid on one plot for comparison.
- **Per-datasource timezone control:** Override the timezone of any supported source via `additional_information.timezone` (all plots still render in the configured display timezone).
- **Export:** Generate standalone HTML visualizations for sharing.

The full list of supported data sources — with folder keywords, accepted file extensions, and typical signals — lives in Section 3: **Patient Data & Supported Data Sources**.

2 Launching the Application

2.1 First launch

ClinicalScope is not code-signed, so every operating system interrupts the **first** launch with a warning about an unknown publisher. This is expected. The steps below get you past it, and they are needed only once — afterwards the application starts normally.

2.1.1 Windows

Unzip the archive, then **right-click ClinicalScope.exe → Run as administrator**. The first launch needs elevation; a plain double-click will not start it.

SmartScreen may also warn that the publisher is unknown — choose **More info → Run anyway**.

This is a one-time step. After that first run, ClinicalScope starts on a normal double-click, including after a reboot.

2.1.2 macOS (Apple Silicon)

macOS quarantines downloaded applications and will not open an unsigned one through the normal flow. Remove the quarantine flag from the **.zip before** unzipping it:

```
cd ~/Downloads
xattr -d com.apple.quarantine ClinicalScope-macOS-arm64.zip
unzip ClinicalScope-macOS-arm64.zip -d ClinicalScope
```

Then run the ClinicalScope executable inside the ClinicalScope folder.

Heads-up: some browsers (Safari in particular) unzip downloads automatically. If yours does, turn that option off or use another browser — the command above must run on the **.zip** file, not on an already-extracted folder. If you have downloaded the archive more than once, check that you are acting on the right **.zip** and the right extraction folder.

There is no Intel Mac build; on those machines, install with `pip install clinical-scope` instead.

2.1.3 Linux

Unzip the archive and run the executable, making it executable first if it does not start:

```
unzip ClinicalScope-linux-x86_64.zip
chmod +x ClinicalScope/ClinicalScope      # only if needed
./ClinicalScope/ClinicalScope
```

2.2 Starting the App

Locate the **ClinicalScope** executable in the application folder and double-click it. The very first time, follow [First launch](#) instead — on Windows the executable has to be started as administrator once.

A terminal window will appear showing the application starting up. After a few seconds, your default web browser will automatically open at:

`http://127.0.0.1:8050`




Figure 1: Application launch screen

The bundle also ships this user guide and a template folder for organizing patient data. A `logs` folder appears after the first run, to help with debugging.

To **close** ClinicalScope, close the terminal window that opened with it — the application runs inside that window. If the window is hidden, end the `ClinicalScope` process from your system’s process manager.

2.3 Application Overview

The interface is organized top-to-bottom in the following order:

1. **Database Options** – Three buttons side-by-side:
 - **Upload config file** (blue) – load a custom `database options` file, either `.json` or `.xlsx`.
 - **Reload last config** (grey) – appears only if a previously uploaded config was cached; restores it with one click.
 - **Default visualization (all sources)** (green) – enables every registered data source with default settings, no file needed.
2. **Patient Options** – Configure data folder, time range, and per-source settings. The form is generated dynamically from the loaded database options.
3. **Action buttons** – **Process visualization** (orange) and **Inspect data** (teal).
4. **Annotations Controls** – Shape dropdown + Modify/Delete buttons (visible after a successful visualization).
5. **Inspection pop-up** – Full-screen overlay triggered by the Inspect button, with per-source status badges, column tables, and a CSV download.
6. **Visualization Area** – Interactive plots.

A **Settings** button sits at the top right, above the Database Options row. It opens your personal display and export settings, which apply to every patient you open — see [Settings](#).




Figure 2: Application main interface

3 Patient Data & Supported Data Sources

Each patient’s data must be organized in a root folder with **one subfolder per data source**. The application automatically identifies data sources based on keywords in subfolder names.

You only need subfolders for the data sources you actually have — empty or missing subfolders are silently skipped. The `clinical_scope_output/` subfolder is created automatically the first time you process a patient. It contains annotations, `.html` visualizations, formatted data.

3.1 Example Folder Layouts

A full setup with several devices:

```
Patient1/
  eit/
  fluxmed_signals/
  fluxmed_parameters/
  servo_u/
  mindray_scope/
  edf/
  other/
  clinical_scope_output/          ← auto-created
```

An “**Other**”-heavy setup — drop any CSV/Parquet file with a datetime column into `other/`; each file is configured independently via `other::<stem>` keys in the database options (see [dedicated section](#) below):

```
Patient1/
  other/
    waves.parquet      → configured under "other::waves"
    numerics.csv       → configured under "other::numerics"
    syringe_log.csv    → configured under "other::syringe_log"
  clinical_scope_output/
```

3.2 Folder Naming Rules

Folder names are **flexible** — they just need to contain the required keywords:

- **Case-insensitive** — `Mindray_Scope`, `MINDRAY-SCOPE`, `mindray scope` all match.
- **Any separator** — underscore, dash, space, or none.
- **Any order** — `scope_mindray` works just as well as `mindray_scope`.
- **Partial keywords don’t match** — `flux` alone will not match `fluxmed_*`; the full word is required.

A few sources are also **identified by file extension** when the folder is ambiguous or generic (e.g., `.asc` files inside a folder are enough to classify it as EIT, `.sta` for Servo-U, `.xml` for Mindray Scope).

3.3 Canonical Data Source Table

Source	Module name	Folder keywords	Accepted extensions (ordered by preference)	Discovery mode	Typical signals
EIT (PulmoVista)	<code>eit</code>	<code>eit</code>	<code>.asc</code>	All files	Global/local impedance, impedance percentages
FluxMed Signals	<code>fluxmed_signals</code>	<code>fluxmed, signals</code>	<code>.parquet, .txt, .csv</code>	Single file	Respiratory waveforms

Source	Module name	Folder keywords	Accepted extensions (ordered by preference)	Discovery mode	Typical signals
FluxMed Parameters	fluxmed_parameters	fluxmed, parameters	.parquet, .txt, .csv	Single file	Respiratory parameters
Servo-U	servo_u	servo	.sta	All files	Ventilator waveforms and settings
Mindray Scope	mindray_scope	mindray	.xml, .csv	All files	Monitor waveforms (ECG, SpO2, pressure)
Mindray Respi Waves	mindray_respi_waves	mindray, resp, wave	.parquet, .csv	Single file	High-frequency respiratory waveforms
Mindray Respi Numerics	mindray_respi_numerics	mindray, resp, numeric	.parquet, .csv	Single file	Respiratory parameters (Vt, RR, PEEP, etc.)
EDF / EDF+	edf	edf	.edf	All files	Any signal an amplifier or recorder exports as EDF (typically EEG)
Other (Generic)	other	other	.parquet, .csv	All files (one entry per file)	Any time-series with a datetime column

Single file sources expect exactly one data file per folder. When several formats coexist (e.g., `data.csv` and `data.parquet`), the most preferred extension wins. If multiple unrelated stems remain after that filter, the source is skipped and a warning is logged.

All files sources load every matching file in the folder and concatenate them. The **Other** source is special: each file produces an independent entry named `other:<stem>` (see below).

EDF recordings. Channels are read exactly as the file stores them — a file holding bipolar derivations (Fp1-F7) shows those, a file holding referential channels shows those. Channels recorded at different sample rates are placed on a shared time axis, each keeping its own sampling.

An EDF header always states a start date and a start time, but de-identification commonly blanks them (the format’s “unknown date” is 01.01.1985), leaving a recording that is effectively relative time only. Place such a recording with the per-source `recording_start` option in `patient_options.json`:

- **A full timestamp** (2024-10-08 08:12:33) puts the first sample at that instant — use this when the file’s clock time was blanked too.
- **A date alone** (2024-10-08) shifts the recording by whole days and keeps the file’s own time of day — use this when only the date was scrubbed.
- Times are read in the **device’s** timezone (the source’s `additional_informations.timezone`, Europe/Paris by default), not your display timezone.

A file that still carries a real start date keeps it, and `recording_start` is ignored. A file with no date and no `recording_start` is still plotted, anchored at 1985-01-01, with a warning in the log.

Per-source configuration options (`field_display`, `signals`, `grouped_fields`, `loop`, `additional_informations`, etc.) are documented in the [Configuration File Reference section](#).




Figure 3: Patient folder structure example

3.4 Generic “Other” Data Source

other is the **generic escape hatch** for any CSV or Parquet file that has a datetime column but does not fit any of the specialised sources — the most natural entry point if your data is already well formatted.

How it works:

- Drop any number of `.csv` or `.parquet` files into an `other/` subfolder; every one is discovered automatically.
- Each file produces an **independent entry** keyed by its stem (filename without extension): `waves.parquet` becomes `other::waves`.
- Columns within a file are exposed as `<stem>::<column_name>` so names stay globally unique (important for cross-datasource groups and loops).

Datetime column auto-detection. The loader automatically finds your file’s datetime column by name and content, including common non-English names — no configuration needed. If no column can be confidently identified as a timestamp, the file is skipped (with a warning) rather than guessed at.

Per-file configuration. In `database_options`, use `other::<stem>` keys — one block per file — exactly like any other datasource:

```
"other::waves": {
  "field_display": ["art", "pleth"],
  "signals": {
    "art": { "label": "Arterial Pressure", "unit": "mmHg", "color": "red" },
    "pleth": { "label": "Plethysmography", "unit": "AU", "color": "purple" }
  },
  "additional_informations": { "timezone": "Europe/Paris" }
},
"other::numerics": {
```

```

    "field_display": ["FC", "SpO2"]
}

```

Per-file plots. A per-file block may also define `grouped_fields`, `loop`, `spectrogram` and `psd`, naming its own columns directly — no `other::<stem>::` prefix needed inside the block:

```

"other::eeg": {
  "spectrogram": {
    "Fp1": { "signal": "Fp1", "freq_range": [0.5, 30.0] }
  }
}

```

The plot is titled with its file’s name in front — the example above appears as `eeg:Fp1` — so two files can each define a PV loop or an Fp1 spectrogram without one replacing the other.

Per-file timezone. Each `other::<stem>` block may declare its own `additional_informations.timezone`, for a file exported by a device in another zone. Without it, timestamps that carry no timezone of their own are read as UTC.

Per-file trace style. A `trace_options` block changes how that file’s traces are drawn — sparse, step-like data (infusion rates, hand-entered values) reads much better with visible points than as a bare line:

```

"other::syringe": {
  "trace_options": { "mode": "lines+markers", "line_width": 2.0 },
  "additional_informations": { "timezone": "Europe/Paris" }
}

```

The same block works in any per-source section; each `other::<stem>` file carries its own, so a curated infusion log and a raw waveform export can look different. Full key list in [trace_options Block](#).

Per-file processing options. Every file you declare with an `other::<stem>` key also gets its own box in the *Specific Options* panel, sitting alongside the device boxes rather than nested under a shared “Other” one. Each box carries that file’s own `time_shift` and *Group signals by source file*, so a curated two-column export and a ninety-column raw dump can be corrected and laid out independently. Files present in the folder but not declared in `database_options` fall back to the shared “Other (generic)” box. See `patient_options.json`.

Inspection shows one entry per file, so you can verify that every file was correctly discovered and parsed — see [Inspect Data](#).

See `example/demo_database/database_options.json` in the source repository for a full reference configuration — its `other::waves`, `other::numerics` and `other::syringe` sections configure the three files shipped in `demo_patient/other/`.

4 Loading Database Options

Database options define **which data sources to enable** and **how signals should be displayed** (labels, units, colors, grouping). There are three ways to get one into the app.

4.1 Option 1: Upload a Custom Configuration File

Click the blue “**Upload config file**” button to load a custom configuration. Two formats are accepted:

- **.json extension** — a JSON file following the structure described in [Configuration File Reference section](#).
- **.xlsx extension** — an Excel spreadsheet following the column layout documented in the same reference. The spreadsheet is converted to the equivalent JSON structure on load.

When the upload succeeds the file is **cached locally** at `~/.clinical_scope/last_database_options.json` — this is what powers Option 2 below. A database options file holds only signal metadata (labels, colors, units, field mappings), never patient data or PHI, so nothing sensitive leaves the patient folder.

4.2 Option 2: Reload Last Config (Daily Workflow)

If a custom configuration was previously uploaded, a grey “**Reload last config**” button appears automatically on startup. Click it to instantly restore the last used configuration without browsing for files — ideal when you re-open the app every day with the same setup.

4.3 Option 3: Default Visualization (Quick Start)

Click the green “**Default visualization (all sources)**” button. This automatically enables **every registered data source** with its built-in default display settings — no configuration file needed. This is the recommended starting point for new users, and it automatically picks up any new data sources added to the library without requiring a config update. Note that if you have hundreds of timeseries, they will all be plot and you may experience performance issues on the app.




Figure 4: Default visualization button

5 Configuring Patient Options

After loading database options, the **Patient Options** form appears, generated dynamically from them.

5.1 Global Options

These apply to all data sources:

Option	Description
Path to data (folder)	Full path to the patient's root data folder
Time start filter	Start of the time window to display (format: YYYY-MM-DD HH:MM:SS). Leave empty to use all available data.
Time end filter	End of the time window to display. Leave empty to use all available data.
Re-use data if already loaded once	When checked, reuses previously cached .parquet files from the <code>clinical_scope_output/</code> folder, significantly speeding up subsequent loads. Un-tick if raw patient data has been modified, or once after updating ClinicalScope

The time filters are typed and shown in the **Display timezone**, set once in [Settings](#) — a label next to the fields names it.

The screenshot shows a web browser window with the URL `http://127.0.0.1:8050`. The page title is "Database Options" and it includes buttons for "Upload config file", "Reload last config", and "Default visualization (all sources)". A green message states "Successfully loaded example_database_options.xlsx". Below this is the "Patient Options" section, which is highlighted with a red box. It contains the "Global Patient Options" section with fields for "Path to data (folder)", "Time start filter", "Time end filter", and a checked checkbox for "Re-use data if already loaded once". Below this is the "Specific Options" section with four cards for different data sources: "Philips scope - waves", "EIT - PulmoVista", "Philips scope - numerics", and "Syringe", each with a "Time shift" field. At the bottom are buttons for "Process visualization" and "Inspect data".

Figure 5: Global patient options

5.2 Per-Source Options

Each data source present in the loaded database options gets its own card below the global options. Common per-source options:

- **Time shift** (seconds): Adjust the time alignment of a source relative to others. Useful when devices were not perfectly synchronized.
- **Day**: Specify the recording date for sources that require it (e.g., EIT data).

The screenshot shows a web browser at `http://127.0.0.1:8050` displaying the 'Database Options' page. The page has a top navigation bar with a search icon, zoom level (90%), and a star icon. Below the navigation bar, there are three buttons: 'Upload config file' (blue), 'Reload last config' (grey), and 'Default visualization (all sources)' (green). A green message bar states 'Successfully loaded example_database_options.xlsx'. The 'API Version: 0.2.0' is displayed in the top right corner.

The main content area is divided into two sections:

- Patient Options**
 - Global Patient Options**
 - Path to data (folder): Path...
 - Time start filter: YYYY-MM-DD HH:MM:SS
 - Time end filter: YYYY-MM-DD HH:MM:SS
 - Re-use data if already loaded once: ☒
- Specific Options** (highlighted with a red border)
 - Philips scope - waves**
 - Time shift: - 0 +
 - Philips scope - numerics**
 - Time shift: - 0 +
 - EIT - PulmoVista**
 - Time shift: - 0 +
 - Day of EIT recording: YYYY-MM-DD HH:MM:SS
 - Syringe**
 - Time shift: - 0 +

At the bottom of the 'Specific Options' section, there are two buttons: 'Process visualization' (orange) and 'Inspect data' (blue).

Figure 6: Per-source options

6 Settings

The **Settings** button at the top right opens your personal settings. They belong to you rather than to a patient or a database: they stay the same whichever data you open, and are saved as soon as you change them, so they survive a restart.

Your settings never overwrite a database options file. They fill the gaps: if the configuration sets a color for a signal, that color is used; if it says nothing, your palette applies. Changes take effect on the next **Process visualization**.

6.1 App behavior

Setting	What it does
Save a full-resolution HTML export on each Process	Writes <code>visualization.html</code> into <code>clinical_scope_output/</code> every time you process. Off by default, because the export takes extra time on large recordings.
Embed Plotly in the HTML export	Makes the exported file self-contained, so it opens on a machine with no internet access. The file grows by roughly 3.5 MB. Leave it off if the file is only ever opened online.
Inspect: read only configured signals	Makes Inspect data read only the signals your database options list, instead of everything in the file — lighter on memory, at the price of no longer seeing the signals you have not configured yet. Off by default (see Inspecting only the signals you configured).

6.2 Plot defaults

Setting	What it does
Display timezone	The timezone every plot and the Patient Options time filters are shown in (IANA name, e.g. Europe/Paris). Changing it does not move your data — the time filters are rewritten to keep pointing at the same instant, just written in the new timezone's local time.
Height of each time-series subplot	In pixels.
Height of each loop subplot	Same, for loop plots — which stay square, so this also sets their width.
Loop subplots per row	How many loops sit side by side, 1 to 3.
Maximum width of one legend entry	Caps how much horizontal room the legend can take from the plot. Longer signal names are truncated.
Palette for signals with no color in the config	Both colorblind-safe palettes are readable under the common color-vision deficiencies.
Plot theme	Light or dark background.
Hover: x-axis time format	Whether the hover panel shows the time only, or the full date and time.
Hover: panel style	<i>Unified</i> lists every trace at the hovered time in one panel; <i>closest point only</i> shows just the nearest trace. Unified is comfortable for a few signals and crowded on a subplot with many.

Setting	What it does
Hover: significant digits of the y value	How precisely hovered values are printed. A signal with its own hover format in the database options keeps it.
Spectrogram colour range — minimum / maximum (dB)	Colour scale bounds used by any spectrogram whose configuration leaves db_range unset.

7 Processing Data

Once patient options are configured, two actions are available from the action row. **Both run the same data pipeline** — find → load → format — and share the same per-datasource progress bar. What differs is only the outcome: the orange button builds interactive plots, the teal button produces a structured summary of what was found.

7.1 Process Visualization (orange button)

Click the large orange “**Process visualization**” button to generate the plots. The steps are:

1. **Validation:** The application verifies that all mandatory fields are filled in and that the data folder exists.
2. **Data Discovery:** For each enabled data source, the application scans the patient folder for matching subfolders and files.
3. **Data Loading:** Raw data files are parsed according to each source’s format.
4. **Formatting:** Signals are filtered, resampled, and converted using your database options (labels, units, time range).
5. **Caching:** Processed data is saved as `.parquet` files in the `clinical_scope_output/` subfolder for faster reloading next time. Tick “**Re-use data if already loaded once**” (`quick_load`) to skip raw parsing on subsequent runs and read the cache instead.
6. **Plot Generation:** Interactive Plotly figures are created and displayed in the visualization area.

While processing runs, a **per-datasource progress bar** appears below the action row. It shows (`completed / total`): `<datasource>` and its color matches the action — orange for visualization. The bar advances as each datasource completes; it never reaches 100% while the last source is still being processed (by design, so you always see which source is active).

A success message appears when processing completes. If no data is found for a source, it is silently skipped.

7.2 Inspect Data (teal button)

The teal “**Inspect data**” button stops before signal extraction and plot building. It is the **recommended first step** when opening a new patient folder: no figures are built, so it is faster, and it immediately exposes loading errors, missing folders, or unexpected column names — before you run the heavier visualization. The same progress bar is shown, colored teal.

When inspection completes, a **full-screen inspection modal** opens, with one section per datasource containing:

- **Status badge** — colored per outcome:
 - green `ok` — data loaded successfully
 - orange `file_not_found` — no matching file/folder in the patient directory
 - red `load_error` or `format_error` — loading or formatting raised an exception (the error message is shown below the badge)
- **File path** — the detected data file or folder.
- **Date ranges** — two lines. *Date range in file* is the source’s own timestamps, untouched. *After time options* is what the application will actually plot, once **every** time setting has been applied: the source’s time shift, its recording start or day for sources that need one, its timezone, and finally the `datetime_start / datetime_end` window. The second line differing from the first is therefore normal and not necessarily a sign that data was cut.
- **Columns table** — each column with: raw name, whether it is configured in the database options, the point count in the file, the count kept after the same time options, and the first and last kept timestamps.

Note for the Other (Generic) source: because the `other` folder may contain several independent files, the inspection modal shows **one entry per file** (e.g. `other::waves`, `other::numerics`) rather than a single aggregated entry. Each entry has its own date range and column list.

A “**Download CSV**” button in the modal header exports the full inspection result as a CSV for offline analysis, sharing, or import into the `generate-database-options` helper.

7.3 Inspecting only the signals you configured

Once your database options are settled, the unconfigured columns are usually noise, and reading them costs memory you may not have on a long recording. The setting “**Inspect: read only configured signals**” ([Settings](#) → App behavior) makes Inspect read only the signals listed in your database options. A section built that way says so above its table, so a column you do not see is one you have not configured — never one missing from the data.

A small “**Configured columns only: on/off**” label sits next to the Inspect data button itself, mirroring the setting so it stays visible even with the Settings modal closed — a safeguard against running a search for a signal you expect to see, on a setting you forgot was on.

Two things stay true whatever you set here:

- **The time window is never applied while reading.** Inspect’s whole point is to compare the file against the window you asked for — how many points survive it, and whether the file even covers it. Reading the window directly would make every source look like a perfect 100% match, and the single most common problem Inspect catches (a window off by an hour, about to plot forty points) would become invisible.
- **The saving only ever applies to a parquet read.** For most sources that means a patient you have already processed once: the first run has to read the manufacturer’s export as it comes, and only the copy the application then saves in `clinical_scope_output/` can be read selectively. On a first-ever inspection of those sources the setting changes nothing, and the table shows every column as usual. The **Other (Generic)** source is the exception — it reads its own CSV/parquet files directly rather than through that cache, so a `.parquet` file there is pruned starting on the very first inspection (a `.csv` file there still shows every column, whatever this setting is).

Database Options

Upload config file | Reload last config | Default visualization (all sources)

Data Inspection Download CSV Close

philips_waves ok

File: /Users/alexis/Codes/ClinicalData/Visualizer/example/example_patients/Patient_full/philips_waves/data_waves_anonymized.parquet
 Date range in file: 04-09-15 08:12:33 → 04-09-15 08:13:28
 After filter: 04-09-15 10:12:38 CEST → 04-09-15 10:13:33 CEST

Column	Configured	Raw pts	Filtered pts	% retained	First (filtered)	Last (filtered)
art	×	10,000	10,000	100.0%	04-09-15 10:12:38 CEST	04-09-15 10:13:33 CEST
vol	×	0	0	—	—	—
pleth	×	10,000	10,000	100.0%	04-09-15 10:12:38 CEST	04-09-15 10:13:33 CEST
paw	×	0	0	—	—	—
p4	×	0	0	—	—	—
debit	×	0	0	—	—	—

eit ok

File: /Users/alexis/Codes/ClinicalData/Visualizer/example/example_patients/Patient_full/cdv_visu/eit.parquet
 Date range in file: 0.5651115972 → 0.5718922685
 After filter: 15-04-09 13:33:45 CEST → 15-04-09 13:43:31 CEST

Column	Configured	Raw pts	Filtered pts	% retained	First (filtered)	Last (filtered)
Image	×	11,718	11,718	100.0%	15-04-09 13:33:45 CEST	15-04-09 13:43:31 CEST
Global	✓	11,718	11,718	100.0%	15-04-09 13:33:45 CEST	15-04-09 13:43:31 CEST
Local 1 (X:11 Y:22)	×	11,718	11,718	100.0%	15-04-09 13:33:45 CEST	15-04-09 13:43:31 CEST
Local 2 (X:21 Y:22)	×	11,718	11,718	100.0%	15-04-09 13:33:45 CEST	15-04-09 13:43:31 CEST
Local 3 (X:11 Y:11)	×	11,718	11,718	100.0%	15-04-09 13:33:45 CEST	15-04-09 13:43:31 CEST

Figure 7: Inspect data pop up

7.4 Inspecting from the Command Line

The same inspection is available as a standalone script:

```
python scripts/inspect_patient_data.py <patient_folder> \  
  [--database-options <path>] \  
  [--patient-options <path>] \  
  [--output-csv out.csv] \  
  [--configured-columns-only] \  
  [--verbose]
```

Without `--database-options`, all registered datasources are inspected with their defaults. Pass `--output-csv` to save the same per-column table the UI download produces. `--configured-columns-only` is the command-line form of the Settings option above, with the same two caveats.

8 Interacting with Plots

8.1 Navigation Controls

Each plot provides a toolbar (top-right corner) with the following tools:

Tool	Action
Zoom	Click and drag to zoom into a rectangular region
Pan	Click and drag to move the view
Zoom In / Zoom Out	Incremental zoom buttons
Autoscale	Reset the view to fit all data
Reset Axes	Return to the original view
Download as PNG	Save the current plot view as an image

Two shortcuts need no toolbar: the **scroll wheel** zooms the x-axis, and a **double-click** resets the axes to show all data.

8.2 Dynamic Resampling (FigureResampler)

For high-frequency signals (e.g., waveforms sampled at hundreds of Hz), the application uses **Plotly-Resampler** to dynamically load detail as you zoom in. A long time range shows a downsampled overview; zooming into a shorter window loads the full-resolution data automatically. This keeps the interface responsive with millions of data points.




Figure 8: Interactive plot navigation

9 Annotations

The annotation system lets you mark events, time windows, or individual points directly on the plots. Annotations are saved to `annotations.json` in the patient folder and persist across sessions.

9.1 Annotation Toolbar

After a successful visualization, the **annotation toolbar** appears above the plots. It contains:

- **Type buttons** — select the annotation type to place next: *Time Event*, *Time Window*, or *Point*.
- **New Group** — create a named group to organise related annotations.
- **Drag** — switch to nudging existing annotations with the mouse instead of placing new ones.
- **Active group display** — shows which group new annotations will belong to.
- **Save** — writes all annotations to disk.
- **Exit mode** — visible whenever a mode is on; click to leave it without saving.

9.2 Annotation Types

Type	Description	Supported plots
Time Event	Vertical line marking a single instant	Time-axis plots only
Time Window	Shaded region spanning a time range	Time-axis plots only
Point	Dot marker at a specific (x, y) location	All plots including loop plots

Click a type button to activate it, then click (or click-and-drag for Time Window) on a plot. A creation modal appears where you can set a **label** and **color** before confirming.

Spectrograms have a time axis like a time-series plot, so all three types can be placed on them. Loop and PSD plots take Point annotations only, because their x-axis is a signal value and a frequency respectively, not a time.

9.3 Moving an Annotation

Each annotation in the list below the toolbar carries a **Move to...** button. Click it, then click the new position on the plot — a Time Window takes two clicks, the new start and then the new end, exactly like placing one. The annotation keeps its label, colour and group; only its position changes. An annotation placed on all subplots stays on all subplots.

Starting a move leaves the active group, so annotations you place afterwards are ungrouped until you click **Continue** on that group again. **Exit mode** cancels a move you have started but not finished.

Use **Move to...** whenever the annotation has to end up somewhere new — on a different subplot, or on a different plot. To correct a mark that is only slightly off, **Drag** below is quicker.

9.4 Dragging an Annotation with the Mouse

Drag turns the annotations already on the plots into handles you can grab: take a Time Window's edge to stretch it, its body to slide it whole while keeping its length, a Time Event to move it, or a Point to reposition it. For a mark that is a few seconds out, this is quicker than **Move to...**

It is a mode you switch on and off, because while it is on the annotations sit over the plots and take the clicks and hover that would otherwise reach the signals underneath — click **Drag** again, or **Exit mode**, when you are done. Placing new annotations and dragging existing ones cannot both be on: turning one on turns the other off.

A drag only ever changes position along the time axis, and a Point's value. Dragging up or down does nothing — release the mouse and the annotation snaps back to where it belongs. Use **Move to...** to put an annotation on a different subplot. One thing to watch: a Point on a loop plot carries the time it was taken from, and a drag cannot work that time out again, so a dragged Point loses it. The annotation list shows each Point's time, so you can see when that has happened.

9.5 Hiding an Annotation

A busy plot — a group of several hundred time windows, say — can be harder to read than the signals underneath it. Two buttons on each annotation in the list control how much of it is drawn:

- **L:on** / **L:off** hides just the text label, leaving the line, rectangle or point marker in place. Useful when you want to see *where* your marks are without reading them.
- **V:on** / **V:off** hides the annotation completely — label, shape and marker alike. Hiding a large group also makes the plot redraw noticeably faster.

Group headers carry the same pair, **Labels: on/off** and **Visible: on/off**, which apply to every annotation in the group at once. If only some members are hidden, the group reads as shown, so one click hides all of it.

While an annotation is hidden its label and **Move to...** buttons are greyed out — their effect would not be visible until you show it again. Hiding never deletes anything: the annotation stays in the list, the count above it reports how many are hidden, and showing it again restores exactly the label setting it had.

Visibility is saved with everything else, so a hidden group stays hidden the next time you open the patient. If you open a file and the plot looks empty, check the annotation count — it will tell you how many are hidden.

9.6 Groups

Annotations can be organised into named groups. Click **New Group**, enter a name, and all subsequent annotations will belong to that group until you switch or create another. A group is not saved as an entity of its own — each annotation carries its group's name, so an empty group does not survive a save.

9.7 Persistence

`annotations.json` sits in the patient data folder, next to the datasource sub-folders. Annotations are reloaded automatically when you re-process the same patient.

If you write the file by hand or generate it from another tool, any extra fields you add to an annotation are kept as they are.

9.8 Python API

Annotations can be loaded programmatically for analysis:

```
from clinical_scope import load_annotations, load_database_annotations

# Single patient - accepts a JSON file, a clinical_scope_output/ folder, or a patient folder
annotations = load_annotations("/data/Patient01")

# Whole database - scans all patient sub-folders and sets ann.patient on each result
all_anns = load_database_annotations("/data")
```

Time shift-0+

Time shift-0+

Process visualization

Inspect data

Visualization succeeded – 2 plot(s) generated.

Annotate:Time EventTime WindowPointNew GroupGroup: typical loop

Exit modeSave14 annotationsSaved (14)

subsequences

Time WindowGlobal (3)ContinueLabels: onDelete all

subsequences

10:29:31 - CO2 - Global

LockX

subsequences

10:30:47 - CO2 - Global

LockX

subsequences

10:37:48 - CO2 - Global

LockX

typical loop

Point (10)ContinueLabels: offDelete all

Other annotations

point1

10:56:30 - CO2 - CO2

LockX

time_series



Figure 9: Annotations tools

10 Configuration File Reference

10.1 patient_options.json

This file defines patient-specific settings. It is automatically saved to the `clinical_scope_output/` subfolder each time you click “Process visualization”.

```
{
  "data_folder": "/path/to/patient/data",
  "datetime_start": "2024-10-08 10:00:00",
  "datetime_end": "2024-10-08 12:00:00",
  "quick_load": false,
  "other::numerics": {
    "time_shift": 20.0
  },
  "eit": {
    "day": "2024-10-08"
  },
  "edf": {
    "recording_start": "2024-10-08 08:12:33"
  },
  "other::waves": {
    "time_shift": 5.0,
    "group_by_file": false
  }
}
```

Key	Type	Default	Description
<code>data_folder</code>	string	—	Path to the patient’s root data folder (required)
<code>output_root</code>	string	""	Writable folder for output when the data folder is read-only. Leave empty to write inside the patient folder. When set, all output goes to <code><output_root>/<patient_folder_name></code>
<code>datetime_start</code>	string or null	null	Start of the time window (YYYY-MM-DD HH:MM:SS). Leave empty to use all available data.
<code>datetime_end</code>	string or null	null	End of the time window. Leave empty to use all available data.
<code>quick_load</code>	boolean	false	Reuse previously cached .parquet files in <code>clinical_scope_output/</code>
<code><source_name></code>	object	—	Per-source options block (e.g., <code>time_shift</code> , <code>day</code>)

Key	Type	Default	Description
<code>other::<stem></code>	object	—	Options for a single file inside <code>other/</code> , taking precedence over the shared <code>other</code> block. See Generic “Other” Data Source .

output_root (read-only data folders). Set this when the patient folder lives on a read-only mount and ClinicalScope cannot write its cache, annotations, or saved configs in place. The layout you already know is reproduced one level deeper, so the root ends up holding one subfolder per patient and `load_database_annotations("<output_root>")` and batch `save_folder` reads work unchanged. Use **one output_root per set of patients** — two patient folders sharing a name (e.g. `patient_01`) under the same root would collide.

10.2 database_options.json

This file controls which data sources are active and how each signal is displayed. A snapshot is automatically saved to `clinical_scope_output/database_options.json` each time you click “Process visualization”.

10.2.1 Top-Level Structure

```
{
  "global": {
    "grouped_fields": { "Pressure": ["other::waves::ART", "servo_u::Airway Press. (cmH2O)] }
  },
  "other::waves": { ... },
  "servo_u": { ... },
  "eit": { ... }
}
```

Each data source key is optional — only include the sources you want to enable. The presence of a source key in this file is what activates that source; removing it disables it entirely.

10.2.2 Per-Source Block Structure

```
"other::waves": {
  "field_display": ["ART", "PAP", "P-aer"],
  "signals": {
    "ART": {
      "label": "Arterial pressure",
      "unit": "mmHg",
      "unit_conversion": 1.0,
      "range": [-10, 200],
      "priority": 1.0,
      "color": "red",
      "visible": true,
      "line_dash": "solid",
      "period_resampling": 0.2,
      "hover_template": null
    }
  },
  "grouped_fields": {
```



```

    "Respiratory": ["P-aer", "CrbVol"]
  },
  "loop": {
    "pv_loop": ["P-aer", "CrbVol"]
  },
  "spectrogram": {
    "Fp1 spectrogram": {
      "signal": "Fp1",
      "freq_range": [0.5, 30.0],
      "db_range": [0, 40]
    }
  },
  "numerics": {
    "period_resampling": 0.5,
    "priority": 1.0
  },
  "trace_options": {
    "mode": "lines+markers",
    "line_width": 2.0
  },
  "additional_informations": {
    "timezone": "Europe/Paris"
  }
}

```

10.2.3 Per-Source Fields Reference

Key	Type	Default	Description
field_display	list of strings	all signals	Signal names to display. Signals absent from this list are loaded but hidden. Omit to show all.
signals	object	{}	Per-signal display options (see Per-Signal Fields Reference below).
grouped_fields	object	{}	Groups of signals to overlay on the same subplot, within this datasource. {"Respiratory waves": ["signal_1", "signal_2", ...], ...}
loop	object	{}	PV-loop definitions: {"loop_name": ["x_signal", "y_signal", ...]}.
spectrogram	object	{}	Spectrogram definitions (see spectrogram below).

Key	Type	Default	Description
<code>psd</code>	object	<code>{}</code>	Power spectral density definitions (see psd below).
<code>numerics</code>	object	<code>{}</code>	Datasource-level defaults applied to every signal (see numerics below).
<code>trace_options</code>	object	source default	How every trace of this datasource is drawn — line, markers, opacity (see trace_options below).
<code>additional_informations</code>	object	<code>{}</code>	Device-level metadata, including timezone override (see additional_informations below).

10.2.4 Per-Signal Fields Reference (`signals.<signal_name>`)

Key	Type	Default	Description
<code>label</code>	string	signal name	Display label shown on plot axes and legends.
<code>unit</code>	string	<code>""</code>	Unit string shown on the Y-axis (e.g., <code>"mmHg"</code> , <code>"cmH2O"</code>).
<code>unit_conversion</code>	float	1.0	Multiplication factor applied to raw values for unit conversion.
<code>range</code>	<code>[min, max]</code> or null	auto	Fixed Y-axis range. Either bound can be <code>null</code> for auto-scaling.
<code>priority</code>	float	source default	Plot ordering priority (lower = higher on page). Signals with the same priority share a subplot.
<code>color</code>	string	auto	Line color (any CSS color string, e.g., <code>"red"</code> , <code>"#1f77b4"</code>).
<code>visible</code>	boolean	<code>true</code>	Set to <code>false</code> to load the signal but hide it from the plot by default.
<code>line_dash</code>	string	<code>"solid"</code>	Line style: <code>"solid"</code> , <code>"dash"</code> , <code>"dot"</code> , <code>"dashdot"</code> .
<code>period_resampling</code>	float	source default	Resampling period in seconds for this specific signal.

Key	Type	Default	Description
hover_template	string	null	Custom hover tooltip. Magic keywords: "fraction" shows values in (0, 1) as 1/n; "percentage" shows them as 33.3%. Any other string is passed directly to Plotly as a <code>hovertemplate</code> . Leave empty for the default compact display.

10.2.5 numerics Block (Datasource-Level Defaults)

The **numerics** block sets **default values** for every signal of a datasource without listing each one — despite the name, it does not select only the numeric ones. Any per-signal entry inside **signals** takes precedence.

```
"other::numerics": {
  "numerics": {
    "period_resampling": 0.5,
    "priority": 2.0
  }
}
```

Key	Type	Default	Description
period_resampling	float	source default	Resampling period in seconds, applied to every signal in this datasource.
priority	float	source default	Default plot priority for numerics (lower = higher on page). Overridden per signal by <code>signals.<name>.priority</code> .

In the Excel format, these values are set via the **sentinel row** (`signal = *`). See [database_options.xlsx](#).

10.2.6 trace_options Block (Datasource-Level Trace Style)

The **trace_options** block changes how every trace of a datasource is drawn. It works in any per-source block — a device datasource, or an `other::<stem>` file scope — and each key you set replaces the datasource's built-in default while the keys you leave out keep theirs.

A dense overlay reads better semi-transparent — the demo database draws the EIT impedance curves this way, so the individual regions stay legible where they cross:

```
"eit": {
  "trace_options": { "opacity": 0.7 }
}
```

Key	Type	Default	Description
mode	string	"lines"	"lines", "markers" or "lines+markers".
line_width	float	source default	Line width in pixels.
line_dash	string	"solid"	Line style: "solid", "dash", "dot", "dashdot".
opacity	float	1.0	Trace opacity, 0–1.
marker_symbol	string	source default	Plotly marker symbol (e.g. "circle", "square"), used when mode includes markers.
marker_size	float	source default	Marker size in pixels.

A key you misspell is reported when the configuration is checked, and ignored. Per-signal `color`, `line_dash` and `visible` live in the `signals` block and win over anything set here — use `trace_options` for the whole datasource and `signals` for the exceptions.

In the Excel format, these values are set via the **sentinel row** (`signal = *`), in the `trace_mode`, `line_width`, `opacity` and `marker_symbol` columns. See [database_options.xlsx](#).

10.2.7 spectrogram Block

The `spectrogram` block defines time-vs-frequency-vs-power plots — one entry produces one subplot, a heatmap with time on the x-axis, frequency on the y-axis, and power as colour. EEG is the main use case, but any sufficiently sampled signal works.

```
"edf": {
  "spectrogram": {
    "Fp1 spectrogram": {
      "signal": "Fp1",
      "freq_range": [0.5, 30.0],
      "db_range": [0, 40]
    }
  }
}
```

Key	Type	Default	Description
signal	string	— (required)	Raw name of the one signal to analyze. No arithmetic — pair or aggregate signals in the source data first if needed.
freq_range	[min_hz, max_hz]	— (required)	Frequency band to display. There is no workable default — pick the band relevant to the signal (e.g. 0.5–30 Hz for EEG).

Key	Type	Default	Description
<code>db_range</code>	<code>[min_db, max_db]</code> or null	your Settings	Fixed colour scale bounds. Left unset, it falls back to your personal colour range setting — fixed rather than auto-scaled, so appearance stays comparable across patients.
<code>window_s</code>	float	derived from <code>freq_range</code>	Advanced: override the analysis window length in seconds. Leave unset unless you know you need it.
<code>overlap</code>	float	0.5	Advanced: override the fraction of overlap between consecutive analysis windows. Leave unset unless you know you need it.

A signal whose `period_resampling` decimates it (see [Per-Signal Fields Reference](#)) cannot be turned into a spectrogram — decimation has no anti-aliasing filter, so the result would show rhythms that are not really there. Remove `period_resampling` for that signal to enable its spectrogram; the log explains which signal and why when this happens.

Power is shown as a spectral density in decibels, so the level a signal reads at does not change with `window_s` or with how fast the signal was sampled. One `db_range` therefore stays meaningful across channels recorded at different rates, and two window lengths of the same channel can be compared directly. Window shape, detrending and averaging follow the defaults of the standard Welch method, so the same numbers come out of any reference implementation.

10.2.8 psd Block

The `psd` block defines power-vs-frequency plots — frequency on the x-axis, power in dB on the y-axis, averaged over the whole loaded time range. Where a spectrogram shows how a rhythm evolves, a PSD shows the shape of the spectrum at rest. One entry produces one subplot, with **one line per signal listed** so several channels can be compared side by side. Typical uses: frequency-domain heart-rate variability, separating ventilation from cardiac components in EIT, and EEG band content.

```
"edf": {
  "psd": {
    "EEG PSD": {
      "signals": ["chan 1", "chan 2", "chan 3"],
      "freq_range": [0.5, 30.0],
      "db_range": [0, 40]
    }
  }
}
```

Key	Type	Default	Description
<code>signals</code>	list	— (required)	Lines to overlay on this plot — see below for the two ways to write one.
<code>freq_range</code>	[min_hz, max_hz]	— (required)	Frequency band to display. There is no workable default — pick the band relevant to the signal (e.g. 0.5–30 Hz for EEG, 0.04–0.4 Hz for heart-rate variability).
<code>db_range</code>	[min_db, max_db] or null	auto	Fixed bounds for the power axis. Left unset, the axis scales to the data.

To compute a PSD over part of a recording rather than all of it, narrow the run with the **Time start** / **Time end** filters — a PSD always covers exactly the time range that was loaded.

Each line takes the colour of its signal, so it matches that signal’s time-series trace. The `period_resampling` restriction described for spectrograms above applies here too, and refusing one signal refuses the whole plot: a comparison missing one of its channels invites a wrong reading more than an absent plot does. The power axis is the same spectral density described for spectrograms above, which is what lets two lines with different `window_s` values sit on one plot.

Writing a `signals` entry. Most of the time a plain name is enough — it is resolved exactly like `grouped_fields` (see [Signal Reference Resolution](#)). Write an object instead when you want to compare the *same* signal analyzed two different ways on the same plot — e.g. a short vs. a long analysis window, to see whether a longer window is smearing two close frequency peaks together:

```
"psd": {
  "EEG PSD": {
    "signals": [
      {"signal": "chan 1", "window_s": 2.0, "label": "narrow window"},
      {"signal": "chan 1", "window_s": 8.0, "label": "wide window"}
    ],
    "freq_range": [0.5, 30.0]
  }
}
```

Key	Type	Default	Description
<code>signal</code>	string	— (required)	The signal for this line, resolved the same way as a plain-string entry.
<code>window_s</code>	float	derived from <code>freq_range</code>	Advanced: override the analysis window length in seconds for this line only.
<code>overlap</code>	float	0.5	Advanced: override the fraction of overlap between consecutive analysis windows for this line only.

Key	Type	Default	Description
<code>label</code>	string	the signal's own name	What this line is called on hover. Required when the same signal appears more than once in one plot — otherwise the two lines are indistinguishable.
<code>color</code>	string	the signal's own colour	Line colour for this line only. Two lines from the same signal otherwise inherit the same colour and overlap indistinguishably.
<code>line_dash</code>	string	the signal's own dash style	<code>solid</code> , <code>dash</code> , <code>dot</code> , <code>dashdot</code> for this line only — a second way (besides colour) to tell two lines from the same signal apart.

10.2.9 additional_informations Block

The `additional_informations` block carries device-level metadata that affects how raw data is interpreted. Currently its only field is `timezone`.

```
"eit": {
  "additional_informations": { "timezone": "UTC" }
}
```

Key	Type	Default	Description
<code>timezone</code>	string	source default	Override the timezone used to give a timezone (e.g., <code>"Europe/Paris"</code> , <code>"UTC"</code>) to timestamps which are timezone-naive. All plots still render in the configured display timezone.

All datasources apply timezone according to the same rule:

- If the loaded data already carries timezone information, it is kept as-is.
- If the data is timezone-naive, the timezone is resolved in this order:
 1. `additional_informations.timezone` in the database options (if present)
 2. The datasource's built-in default timezone

In the Excel format, set the timezone via the **sentinel row** (`signal = *`, `timezone` column). See [database_options.xlsx](#).

10.2.10 Global Fields

```
"global": {
  "grouped_fields": {
    "Pressure": ["ART", "PNIs", "PNIm", "PNId"]
  },
  "loop": {
    "PV loop (ventilator vs. FluxMed)": [
      "servo_u::Paw", "fluxmed_signals::Vt"
    ]
  }
}
```

Key	Description
<code>global.grouped_fields</code>	Groups signals from different datasources onto the same subplot. Signal names must be unique across the datasources involved — use <code>datasource::raw_name</code> to disambiguate.
<code>global.loop</code>	Cross-datasource phase loops: <code>{"loop_name": ["x_ref", "y_ref"]}</code> . Both signals can come from different datasources. Each <code>*_ref</code> is resolved by the same three-mode chain as <code>grouped_fields</code> .

10.2.11 Signal Reference Resolution

Signal references in `grouped_fields`, `loop` and `psd` are resolved by the following three-mode lookup chain — in `global.grouped_fields` and `global.loop`, and in the per-source blocks alike:

1. **Qualified reference** `datasource::raw_name` — explicit and unambiguous. Recommended when the same column name exists in several datasources.
2. **Display name** — the `label` from the signals block. Works when the label is unique across sources; a warning is logged if it matches several signals.
3. **Raw name fallback** — the column name as it appears in the raw data file.

A signal from a file inside `other/` can be written either way: `other::waves::art` is the fully qualified form, and `waves::art` is the raw-name form — both reach the same signal.

One consequence: if a file inside `other/` is named after a data source — `other/servo_u.parquet` — then `servo_u::Paw` could mean either that file's `Paw` column or the real servo-u one. The data source always wins, and the log tells you so and gives you the fully qualified spelling (`other::servo_u::Paw`) that reaches the file's column instead.

Multi-cycle loops are rendered with a **time-range slider** below the plot so you can scroll through cycles.


```

{} example_database_options_other.json M X
example > option_files > {} example_database_options_other.json > {} other::waves > {} additional_informations

2  "global": {
3    "grouped_fields": {
4      "ART pressure (waves + numerics)": [
5        "other::waves::art",
6        "other::numerics::ARTm"
7      ]
8    },
9    "loop": {
10     "Other numerics vs Other mindray_respi_waves": [
11       "other::numerics::ARTm",
12       "other::mindray_respi_waves::MDC_PRESS_AWAY-MDC_DIM_CM_H20"
13     ]
14   }
15 },
16
17 "other::waves": {
18   "field_display": ["art", "pleth"],
19   "additional_informations": {
20     "timezone": "Europe/Paris"
21   },
22   "signals": {
23     "art": {
24       "label": "Arterial Pressure",
25       "unit": "mmHg",
26       "color": "red",
27       "range": [0, 500],
28       "hover_template": "<b>ART:</b> %{y:.1f} mmHg<extra></extra>"
29     },
30     "pleth": {
31       "label": "Plethysmography",
32       "unit": "AU",
33       "color": "purple"
34     }
35   },
36   "grouped_fields": {
37     "Hemodynamic waves": ["art", "pleth"]
38   },
39   "loop": {
40     "Art vs Pleth (within-file)": ["art", "pleth"]
41   }
42 },
43
44
45 "syringe": {
46   "field_display": [
47     "Vitesse Propofol",
48     "Vitesse Remifentanyl"

```

Figure 10: example json database options file

10.3 database_options.xlsx

For clinical users, the **Excel format is the recommended way** to configure signals — it requires no knowledge of JSON and can be edited in any spreadsheet application. On upload, it is automatically converted to the equivalent [JSON structure](#), so every option available in JSON is also available in the spreadsheet.

The file must contain a sheet named **signals**, and may add **loops**, **spectrograms** and **psds**. An optional sheet that is absent or malformed is silently skipped.

10.3.1 signals sheet

One row per signal. The columns **datasource** and **signal** are mandatory; all others are optional and fall back to the defaults listed in the per-signal table above.

Use ***** in the **signal** column to write a **sentinel row** that sets datasource-level defaults (e.g., a common **period_resampling** or **timezone**) without defining a specific signal — equivalent to the [numerics](#) and [additional_informations blocks](#) in JSON. Column names are case-insensitive (e.g., **Label**, **UNIT**, **Hover_Template** all work).

The **Scope** column below indicates where each field is meaningful:

- **Both** — valid in sentinel (*) and per-signal rows
- **Signal** — per-signal rows only; ignored in sentinel rows
- **Sentinel** — sentinel (*) rows only; a warning is logged if set in a per-signal row

Column	Required	Scope	Description
datasource	Yes	Both	Data source name (e.g., servo_u , eit , other::waves).
signal	Yes	Both	Raw signal name. Use * for a sentinel row that sets datasource-level defaults.
label	No	Signal	Display label. Defaults to the signal name if empty or identical.
unit	No	Signal	Unit string (e.g., mmHg).
unit_conversion	No	Signal	Numeric multiplier for unit conversion.
range_min	No	Signal	Minimum Y-axis value.
range_max	No	Signal	Maximum Y-axis value.
priority	No	Both	Plot priority (float). In a sentinel row sets the datasource-level default; in a signal row overrides it for that signal only.
color	No	Signal	CSS color string.
visible	No	Signal	yes / no (default: yes). Set no to draw the signal but start it hidden — it stays in the legend, and one click brings it back. Accepts yes , 1 , true , oui , vrai (case-insensitive).

Column	Required	Scope	Description
<code>line_dash</code>	No	Signal	<code>solid</code> , <code>dash</code> , <code>dot</code> , <code>dashdot</code> .
<code>period_resampling</code>	No	Both	Resampling period in seconds. In a sentinel row sets the datasource-level default; in a signal row overrides it for that signal only.
<code>hover_template</code>	No	Signal	Hover tooltip format. Magic keywords: <code>"fraction"</code> shows values in (0, 1) as 1/n; <code>"percentage"</code> shows them as 33.3%. Any other string is forwarded directly to Plotly as a <code>hovertemplate</code> .
<code>display</code>	No	Signal	<code>yes</code> / <code>no</code> — whether to add this signal to the display list. Default: <code>yes</code> . Set <code>no</code> to keep the row's label and unit on file while leaving the signal out of the plots entirely: unlike <code>visible</code> , it produces no trace and no legend entry. Useful for parking a signal you may want back later, or for describing a column that is not worth plotting (see <code>Comments(-)</code> under <code>fluxmed_parameters</code> in <code>example/demo_database/database_opt</code>).
<code>groups</code>	No	Signal	Semicolon-separated group names (e.g., <code>Respiratory;Pressure</code>). Groups within one datasource become local <code>grouped_fields</code> ; groups spanning multiple datasources become <code>global.grouped_fields</code> .

Column	Required	Scope	Description
<code>timezone</code>	No	Sentinel	Override the timezone for this datasource (e.g., "Europe/Paris", "UTC"). Only valid in * rows; a warning is logged if placed in a per-signal row. Works with <code>other::<stem></code> datasource keys. See additional_informations Block for which datasources support this.
<code>trace_mode</code>	No	Sentinel	<code>lines</code> , <code>markers</code> or <code>lines+markers</code> for every trace in this datasource, <code>other::<stem></code> keys included — e.g. a sparse infusion log reads better as <code>markers</code> than as connected <code>lines</code> . Only valid in * rows; see trace_options .
<code>line_width</code>	No	Sentinel	Line width in pixels for every trace in this datasource. Only valid in * rows.
<code>opacity</code>	No	Sentinel	Trace opacity, 0-1. Only valid in * rows.
<code>marker_symbol</code>	No	Sentinel	Plotly marker symbol (e.g., <code>circle</code> , <code>square</code>) used when <code>trace_mode</code> includes <code>markers</code> . Only valid in * rows.

10.3.2 loops sheet (optional)

One row per PV-loop definition — equivalent to the `loop` key in the [JSON per-source block](#) or in `global`.

Column	Required	Description
<code>datasource</code>	Yes	Data source that owns both signals.
<code>loop_name</code>	Yes	Name for the loop plot (e.g., <code>pv_loop</code>).
<code>x_signal</code>	Yes	Signal name for the X axis.
<code>y_signal</code>	Yes	Signal name for the Y axis.

10.3.3 spectrograms sheet (optional)

One row per spectrogram definition — equivalent to the `spectrogram` key in the [JSON per-source block](#).

Column	Required	Description
<code>datasource</code>	Yes	Data source that owns the signal.
<code>spectrogram_name</code>	Yes	Name for the spectrogram plot (e.g., <code>Fp1 spectrogram</code>).
<code>signal</code>	Yes	Raw name of the signal to analyze.
<code>freq_min</code>	Yes	Minimum frequency to display (Hz).
<code>freq_max</code>	Yes	Maximum frequency to display (Hz).
<code>db_min</code>	No	Colour scale minimum (dB). Leave both <code>db_min</code> and <code>db_max</code> empty to use your personal colour range setting.
<code>db_max</code>	No	Colour scale maximum (dB). Must be set together with <code>db_min</code> , or both are ignored.
<code>window_s</code>	No	Advanced: override the analysis window length in seconds. Leave unset unless you know you need it.
<code>overlap</code>	No	Advanced: override the fraction of overlap between consecutive analysis windows. Leave unset unless you know you need it.

10.3.4 psds sheet (optional)

One row per signal — equivalent to the `psd` key in the [JSON per-source block](#). A signal can belong to zero, one, or several PSD plots, just like the `groups` column on the [signals sheet](#). Signals sharing a group are overlaid on the same plot, one line each.

Column	Required	Description
<code>datasource</code>	Yes	Data source that owns the signal.
<code>groups</code>	Yes	Semicolon-separated PSD plot names (e.g., <code>EEG PSD;Low band</code>). Each name becomes its own plot; signals sharing a name overlay on it. Leave empty to exclude the signal from every PSD plot.
<code>signal</code>	Yes	Raw name of the signal for this line.
<code>freq_min</code>	Yes	Minimum frequency to display (Hz). Read from the first row that introduces each plot; a later row with a different value logs a warning and is ignored.
<code>freq_max</code>	Yes	Maximum frequency to display (Hz). Same first-row-wins rule as <code>freq_min</code> .

Column	Required	Description
<code>db_min</code>	No	Power axis minimum (dB). Leave both <code>db_min</code> and <code>db_max</code> empty to scale the axis to the data. Same first-row-wins rule as <code>freq_min</code> .
<code>db_max</code>	No	Power axis maximum (dB). Must be set together with <code>db_min</code> , or both are ignored.
<code>window_s</code>	No	Advanced: override the analysis window length in seconds for this row's line only. Leave unset unless you know you need it.
<code>overlap</code>	No	Advanced: override the fraction of overlap between analysis windows for this row's line only.
<code>label</code>	No	What this line is called on hover. Required when the same signal appears more than once in one plot (e.g. comparing two <code>window_s</code> values) — otherwise the two lines are indistinguishable.
<code>color</code>	No	Line colour for this row's line only. Leave unset to inherit the signal's own colour — two rows for the same signal then need this to tell their lines apart.
<code>line_dash</code>	No	<code>solid</code> , <code>dash</code> , <code>dot</code> , <code>dashdot</code> for this row's line only. A second way (besides colour) to tell two lines from the same signal apart.

See `example/demo_database/` in the source repository for a complete example in both formats — `database_options.xlsx` (all four sheets) and `database_options.json`, its exact equivalent. Both configure the shipped `demo_patient/`, so you can load either one and press Process straight away.

11 Troubleshooting

11.1 Browser Does Not Open Automatically

If the browser does not open after launching the application, manually navigate to:

`http://127.0.0.1:8050`

Ensure no other application is using port 8050 (typically a previous app launch terminal tab not yet closed). If needed, close the terminal window which was opened in the app and restart the application.

11.2 No Data Found

If the visualization is empty or a data source shows no signals:

- Start by inspecting the data (teal button) rather than visualizing it — it reports much more about what the app could gather from the data.
- Try the default visualization database options, to rule out your own options hiding the data.
- Verify that the **data folder path** is correct and accessible.
- Check that subfolders follow the **naming conventions** (see Section 3).
- Ensure the subfolder contains files with one of the **accepted extensions** for that data source (see Section 3). Files with unrecognized extensions are silently ignored.
- For **single-file** sources: if the folder contains multiple unrelated data files (different stems), the source is skipped. Keep only one data file per folder, or provide the same data in multiple formats (e.g., `data.csv` + `data.parquet`) and the preferred format will be selected automatically.

11.3 Slow Loading

Large datasets may take time to load on the first run. To speed up subsequent loads:

- Enable the “**Re-use data if already loaded once**” (`quick_load`) option. This uses the cached `.parquet` files in `clinical_scope_output/` instead of re-reading raw data files.

11.4 Time Alignment Issues

If signals from different sources appear misaligned in time:

- Use the **Time shift** option in the per-source settings to adjust alignment.
- Verify that the correct **day** or **date** is set for sources that require it (e.g., EIT).

11.5 Application Crashes or Errors

- Check the terminal window for error messages.
- Log files are available in the `logs/` directory.
- Ensure the data files are not corrupted or truncated.
- If you think you are facing a real bug, please report it on the [GitHub issues page](#).

11.6 Known limitations

These may be tackled in the future, depending on what users need — ask for any of them on the [issues page](#).

- No timeshift inside a datasource, e.g. if 2 timeseries from `servo_u` are not aligned, this currently can’t be solved in the app. (Files inside `other/` are the exception — each gets its own `time_shift`.)
- `output_root` keys each patient by its folder name only, so two patient folders with the same name under one `output_root` overwrite each other — see `patient_options.json`.
- EIT recordings spanning more than one calendar day are unsupported: the device’s own files carry no date, so the day is inferred from the **day** option (or, if unset, from **Time start filter**) and applied to the whole recording.

- Spectrograms and PSDs expose only `freq_range`, `db_range`, `window_s` and `overlap`. Everything else about the analysis — window shape, detrending, how windows are averaged, how recording gaps and irregular timestamps are handled — is fixed at the standard Welch defaults and cannot be changed from the config.
- A spectrogram or PSD figure does not state the settings it was drawn with, nor whether anything happened to the data on the way (timestamps regularised, gaps skipped). Keep the `database_options` file alongside the figure if you need that record — for instance to describe the analysis in a publication.